

01_preprocessing

December 21, 2019

1 Kaggle Titanic survival - data preprocessing

Can we predict which passengers would survive the sinking of the Titanic?

Original kaggle page:<https://www.kaggle.com/c/titanic>

Subsequent machine learning notebooks using Titanic survival also provide links to load preprocessed data directly, so this notebook is not strictly needed before using other notebooks, but processing data into a useable form is often a key stage of any machine learning project, and so all practitioners will want to get to grips with common methods.

This Nntebook introduces the following:

- Using Pandas to load and process data (though some familiarity with Pandas is assumed)
- Looking at data types
- Listing feature headings
- Showing data
- Showing a statistical summary of data
- Filling in (imputing) missing data
- Encoding non-numerical fields
- Removing unwanted columns
- Saving processed data

The data includes.

Variable	Definition
survival	Survival (0 = No, 1 = Yes)
pclass	Ticket class
sex	Sex
Age	Age in years
sibsp	# of siblings / spouses aboard the Titanic
parch	# of parents / children aboard the Titanic
ticket	Ticket number
fare	Passenger fare
cabin	Cabin number
embarked	Port of Embarkation(C=Cherbourg, Q=Queenstown, S=Southampton)

1.1 Load modules

```
[1]: import pandas as pd
import numpy as np
```

2 Load data

Data should be in a sub folder named data.

It may be downloaded from:

https://gitlab.com/michaelallen1966/1908_coding_club_kaggle_titanic/tree/master/data

Usually the first thing we will do is split data in training and test (usually with randomisation first), and we hold back the test data until model building is complete. In the case of this kaggle data a separate test data set is supplied, so we do not need to hold back and of the data.

We will load the kaggle data and make a copy we will work on (so we can always refer back to the original data if we wish).

```
[2]: download_required = True

if download_required:

    # Download processed data:
    address = 'https://raw.githubusercontent.com/MichaelAllen1966/' + \
              '1804_python_healthcare/master/titanic/data/train.csv'

    data = pd.read_csv(address)

    # Create a data subfolder if one does not already exist
    import os
    data_directory = './data/'
    if not os.path.exists(data_directory):
        os.makedirs(data_directory)

    # Save data
    data.to_csv(data_directory+'train.csv')
```

```
[3]: original_data = pd.read_csv('./data/train.csv')
data = original_data.copy()
```

Let's have a look at some general information on the table.

```
[4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
```

```
Data columns (total 14 columns):
Unnamed: 0      891 non-null int64
Unnamed: 0.1    891 non-null int64
PassengerId     891 non-null int64
Survived        891 non-null int64
Pclass          891 non-null int64
Name            891 non-null object
Sex             891 non-null object
Age            714 non-null float64
SibSp           891 non-null int64
Parch           891 non-null int64
Ticket          891 non-null object
Fare            891 non-null float64
Cabin           204 non-null object
Embarked        889 non-null object
dtypes: float64(2), int64(7), object(5)
memory usage: 97.6+ KB
```

At this point we can note we have 891 passengers, but that 'Age', 'Cabin' and 'Embarked' have some data missing.

Let's list the data fields:

```
[5]: list(data)
```

```
[5]: ['Unnamed: 0',
      'Unnamed: 0.1',
      'PassengerId',
      'Survived',
      'Pclass',
      'Name',
      'Sex',
      'Age',
      'SibSp',
      'Parch',
      'Ticket',
      'Fare',
      'Cabin',
      'Embarked']
```

Let's look at the top of our data.

```
[6]: data.head()
```

```
[6]:   Unnamed: 0  Unnamed: 0.1  PassengerId  Survived  Pclass  \
0           0           0           1           0         3
1           1           1           2           1         1
2           2           2           3           1         3
3           3           3           4           1         1
```

4	4	4	5	0	3
---	---	---	---	---	---

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

We can count the number of empty values. We can see that we will need to deal with 'age', 'cabin', and 'embarked'.

```
[7]: data.isna().sum()
```

```
[7]: Unnamed: 0      0
      Unnamed: 0.1    0
      PassengerId    0
      Survived       0
      Pclass         0
      Name           0
      Sex            0
      Age           177
      SibSp          0
      Parch          0
      Ticket         0
      Fare           0
      Cabin         687
      Embarked       2
      dtype: int64
```

2.1 Showing a summary of the data

We can use the pandas `describe()` method to show a summary of the data. Note that this only shows numerical data.

```
[8]: data.describe()
```

```
[8]:      Unnamed: 0  Unnamed: 0.1  PassengerId  Survived  Pclass  \
count  891.000000    891.000000    891.000000    891.000000  891.000000
mean   445.000000    445.000000    446.000000     0.383838    2.308642
```

std	257.353842	257.353842	257.353842	0.486592	0.836071
min	0.000000	0.000000	1.000000	0.000000	1.000000
25%	222.500000	222.500000	223.500000	0.000000	2.000000
50%	445.000000	445.000000	446.000000	0.000000	3.000000
75%	667.500000	667.500000	668.500000	1.000000	3.000000
max	890.000000	890.000000	891.000000	1.000000	3.000000

	Age	SibSp	Parch	Fare
count	714.000000	891.000000	891.000000	891.000000
mean	29.699118	0.523008	0.381594	32.204208
std	14.526497	1.102743	0.806057	49.693429
min	0.420000	0.000000	0.000000	0.000000
25%	20.125000	0.000000	0.000000	7.910400
50%	28.000000	0.000000	0.000000	14.454200
75%	38.000000	1.000000	0.000000	31.000000
max	80.000000	8.000000	6.000000	512.329200

Of most likely useful fields we are missing sex and whether a patient embarked or not. So let's code those numerically.

2.2 Filling in (imputing) missing data

For numerical data we may commonly choose to impute missing values with zero, mean or median. We will use the median for age.

We will also create a new column showing which values were imputed (this may be useful information in a machine learning model)

```
[9]: def impute_missing_with_median(_series):
    """
    Replace missing values in a Pandas series with median,
    Returns a completed series, and a series showing which values are imputed
    """
    # Copy the series to avoid change to the original series.
    series = _series.copy()
    median = series.median()
    missing = series.isna()
    series[missing] = median

    return series, missing
```

```
[10]: age, imputed = impute_missing_with_median(data['Age'])
data['Age'] = age
data['AgeImputed'] = imputed
```

We will impute missing embarked text with a 'missing' label

```
[11]: def impute_missing_with_missing_label(_series):
        """Replace missing values in a Pandas series with the text 'missing'"""
        # Copy the series to avoid change to the original series.
        series = _series.copy()
        missing = series.isna()
        series[missing] = 'missing'

        return series, missing

[12]: embarked, imputed = impute_missing_with_missing_label(data['Embarked'])
data['Embarked'] = embarked
data['EmbarkedImputed'] = imputed
```

3 Sorting out cabin data

Cabin data is messy! Some passengers have more than one cabin (in which case we will split out the multiple cabins and just use the first one). Cabin numbers are a letter followed by a number. We will separate out the letter and the number.

```
[13]: # Get cabin data from dataframe
cabin = data['Cabin']

# Set up strings to add each passenger data to
CabinLetter = []
CabinLetterImputed = []
CabinNumber = []
CabinNumberImputed = []

# Convert all cabin data to string (empty cells are current stored as 'float')
cabin = cabin.astype(str)

# Iterate through rows
for index, value in cabin.items():
    # If cabin info is missing (string is 'nan' then add imputed data)
    if value == 'nan':
        CabinLetter.append('missing')
        CabinLetterImputed.append(True)
        CabinNumber.append(0)
        CabinNumberImputed.append(True)
    # Otherwise split string by spaces where there are multiple cabins
    else:
        # Split multiple cabins
        cabins = value.split(' ')
        # Take first cabin
        use_cabin = cabins[0]
```

```

letter = use_cabin[0] # First letter
CabinLetter.append(letter)
CabinLetterImputed.append(False)
if len(use_cabin) > 1:
    number = use_cabin[1:]
    CabinNumber.append(number)
    CabinNumberImputed.append(False)
else:
    CabinNumber.append(0)
    CabinNumberImputed.append(True)

data['CabinLetter'] = CabinLetter
data['CabinLetterImputed'] = CabinLetterImputed
data['CabinNumber'] = CabinNumber
data['CabinNumberImputed'] = CabinNumberImputed

data.drop('Cabin', axis=1, inplace=True)

```

```
[14]: data.head()
```

```

[14]:   Unnamed: 0  Unnamed: 0.1  PassengerId  Survived  Pclass  \
0          0          0          1          0          3
1          1          1          2          1          1
2          2          2          3          1          3
3          3          3          4          1          1
4          4          4          5          0          3

                                     Name      Sex  Age  SibSp  \
0                        Braund, Mr. Owen Harris  male  22.0    1
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0    1
2                        Heikkinen, Miss. Laina  female  26.0    0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)  female  35.0    1
4                        Allen, Mr. William Henry  male  35.0    0

   Parch      Ticket    Fare Embarked  AgeImputed  EmbarkedImputed  \
0      0      A/5 21171   7.2500         S        False            False
1      0      PC 17599  71.2833         C        False            False
2      0  STON/O2. 3101282   7.9250         S        False            False
3      0      113803  53.1000         S        False            False
4      0      373450   8.0500         S        False            False

   CabinLetter  CabinLetterImputed  CabinNumber  CabinNumberImputed
0      missing                True           0                True
1           C                False          85                False
2      missing                True           0                True
3           C                False         123                False

```

4	missing	True	0	True
---	---------	------	---	------

Let's check our missing numbers totals again

```
[15]: data.isna().sum()
```

```
[15]: Unnamed: 0          0
      Unnamed: 0.1      0
      PassengerId      0
      Survived         0
      Pclass           0
      Name             0
      Sex              0
      Age              0
      SibSp            0
      Parch            0
      Ticket           0
      Fare             0
      Embarked         0
      AgeImputed       0
      EmbarkedImputed  0
      CabinLetter      0
      CabinLetterImputed 0
      CabinNumber      0
      CabinNumberImputed 0
      dtype: int64
```

3.1 Encoding non-numerical fields.

There are three types of non-numerical field:

- Dichotomous, which have two, and only two, possibilities (e.g. male/female, alive/dead). These may be recoded as 0 or 1.
- Categorical, which have any number of possibilities that cannot be ordered in any sensible way (e.g. colour of car). Each possibility is coded separately as 0/1 (e.g red = 0 or 1, green = 0 or 1, blue = 0 or 1). This is called 'one-hot encoding' as there will be one '1' (hot) in a set of columns (with all other values being zero).
- Ordinal, which have any number of possibilities but which may be ordered in a sensible way and coded by order of list. For example the size of shirts may be xs, s, m, l and xl. These may be re-coded as size 0, 1, 2, 3, 4 (or scaled in another way if appropriate).

We'll look at sex first. Let's pull that out as a separate 'series'

```
[16]: sex = data['Sex']
      sex.head()
```



```
[16]: 0    male
      1   female
      2   female
      3   female
      4    male
      Name: Sex, dtype: object
```

From looking at the data it appears passengers are either male or female, but data can contain missing values or spelling mistakes, so let's check all the values present. An easy way to do this is to use Python's `set` command which only allows one instance of each value.

```
[17]: set(sex)
```

```
[17]: {'female', 'male'}
```

That's good. We have just 'female' and 'male'. Let's code a new 'male' column manually, and check the mean (the proportion of passengers who are male).

```
[18]: male = data['Sex'] == 'male'
      male.mean()
```

```
[18]: 0.6475869809203143
```

That looks reasonable. We'll add our new column to our dataframe, and remove the old 'sex' column.

To remove a column we use the pandas `drop()` method. To show it is a column we specify `axis=1`. To instruct removal from the data itself we use `inplace=True`. This is the equivalent of saying `data = data.drop()`.

```
[19]: data['male'] = male
      data.drop(['Sex'], axis=1, inplace=True)
```

Let's look at our table now.

```
[20]: data.head()
```

```
[20]:  Unnamed: 0  Unnamed: 0.1  PassengerId  Survived  Pclass  \
0           0           0           1           0         3
1           1           1           2           1         1
2           2           2           3           1         3
3           3           3           4           1         1
4           4           4           5           0         3

      Name  Age  SibSp  Parch  \
0  Braund, Mr. Owen Harris  22.0    1     0
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  38.0    1     0
2  Heikkinen, Miss. Laina  26.0    0     0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)  35.0    1     0
```

4		Allen, Mr. William Henry	35.0	0	0
---	--	--------------------------	------	---	---

	Ticket	Fare	Embarked	AgeImputed	EmbarkedImputed	\
0	A/5 21171	7.2500	S	False	False	
1	PC 17599	71.2833	C	False	False	
2	STON/O2. 3101282	7.9250	S	False	False	
3	113803	53.1000	S	False	False	
4	373450	8.0500	S	False	False	

	CabinLetter	CabinLetterImputed	CabinNumber	CabinNumberImputed	male
0	missing	True	0	True	True
1	C	False	85	False	False
2	missing	True	0	True	False
3	C	False	123	False	False
4	missing	True	0	True	True

Let's do the same with 'embarked'.

```
[21]: embarked = data['Embarked']
      set(embarked)
```

```
[21]: {'C', 'Q', 'S', 'missing'}
```

Ah, we have four possibilities!

We could frame this as a series of if/elif/else statements. That is reasonable for a few possibilities, but what if we have many? We could write our own function to 'one-hot' encode this column, but pandas can already do this for us with the `get_dummies` method.

Note that we pass a couple of useful arguments: `prefix` allows us to add some text to each label, and `dummy_na=True` allows us to specifically code missing values (though we have already given them the label 'missing').

As ever, it is often useful to look at the help for these methods (`help(pd.get_dummies)`).

```
[22]: embarked_coded = pd.get_dummies(embarked, prefix='Embarked')
      embarked_coded.head()
```

```
[22]:   Embarked_C  Embarked_Q  Embarked_S  Embarked_missing
0           0           0           1           0
1           1           0           0           0
2           0           0           1           0
3           0           0           1           0
4           0           0           1           0
```

Nice! We'll add our new table to the data table and drop the original 'Embarked' column. Pandas `concat` method will join our dataframes.

Pandas has `concat`, `merge` and `join` methods for combining dataframes https://pandas.pydata.org/pandas-docs/stable/user_guide/merging.html

```
[23]: data = pd.concat([data, embarked_coded], axis=1)
data.drop(['Embarked'], axis=1, inplace=True)
data.head()
```

```
[23]:
```

	Unnamed: 0	Unnamed: 0.1	PassengerId	Survived	Pclass	\
0	0	0	1	0	3	
1	1	1	2	1	1	
2	2	2	3	1	3	
3	3	3	4	1	1	
4	4	4	5	0	3	

	Name	Age	SibSp	Parch	\
0	Braund, Mr. Owen Harris	22.0	1	0	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	38.0	1	0	
2	Heikkinen, Miss. Laina	26.0	0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	35.0	1	0	
4	Allen, Mr. William Henry	35.0	0	0	

	Ticket	...	EmbarkedImputed	CabinLetter	CabinLetterImputed	\
0	A/5 21171	...	False	missing	True	
1	PC 17599	...	False	C	False	
2	STON/O2. 3101282	...	False	missing	True	
3	113803	...	False	C	False	
4	373450	...	False	missing	True	

	CabinNumber	CabinNumberImputed	male	Embarked_C	Embarked_Q	Embarked_S	\
0	0	True	True	0	0	1	
1	85	False	False	1	0	0	
2	0	True	False	0	0	1	
3	123	False	False	0	0	1	
4	0	True	True	0	0	1	

	Embarked_missing
0	0
1	0
2	0
3	0
4	0

[5 rows x 22 columns]

```
[24]: cabin_coded = pd.get_dummies(CabinLetter, prefix='CabinLetter')
cabin_coded.head()
```

```
[24]:
```

	CabinLetter_A	CabinLetter_B	CabinLetter_C	CabinLetter_D	CabinLetter_E	\
0	0	0	0	0	0	
1	0	0	1	0	0	

2	0	0	0	0	0
3	0	0	1	0	0
4	0	0	0	0	0

	CabinLetter_F	CabinLetter_G	CabinLetter_T	CabinLetter_missing
0	0	0	0	1
1	0	0	0	0
2	0	0	0	1
3	0	0	0	0
4	0	0	0	1

Now let's add those back to the table

```
[25]: data = pd.concat([data, cabin_coded], axis=1)
data.drop(['CabinLetter'], axis=1, inplace=True)
```

```
[26]: data.head()
```

```
[26]: Unnamed: 0  Unnamed: 0.1  PassengerId  Survived  Pclass  \
0          0          0          1          0          3
1          1          1          2          1          1
2          2          2          3          1          3
3          3          3          4          1          1
4          4          4          5          0          3
```

	Name	Age	SibSp	Parch	\
0	Braund, Mr. Owen Harris	22.0	1	0	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	38.0	1	0	
2	Heikkinen, Miss. Laina	26.0	0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	35.0	1	0	
4	Allen, Mr. William Henry	35.0	0	0	

	Ticket	...	Embarked_missing	CabinLetter_A	CabinLetter_B	\
0	A/5 21171	...	0	0	0	
1	PC 17599	...	0	0	0	
2	STON/O2. 3101282	...	0	0	0	
3	113803	...	0	0	0	
4	373450	...	0	0	0	

	CabinLetter_C	CabinLetter_D	CabinLetter_E	CabinLetter_F	CabinLetter_G	\
0	0	0	0	0	0	
1	1	0	0	0	0	
2	0	0	0	0	0	
3	1	0	0	0	0	
4	0	0	0	0	0	

	CabinLetter_T	CabinLetter_missing
--	---------------	---------------------

0	0	1
1	0	0
2	0	1
3	0	0
4	0	1

[5 rows x 30 columns]

Now we will drop the Name and Ticket column (they may perhaps be useful in some way, but we'll simplify things by removing them)

3.2 Drop columns

```
[27]: cols_to_drop = ['Name', 'Ticket']
data.drop(cols_to_drop, axis=1, inplace=True)
data.head()
```

```
[27]: Unnamed: 0  Unnamed: 0.1  PassengerId  Survived  Pclass   Age  SibSp  \
0          0          0          1          0        3  22.0    1
1          1          1          2          1        1  38.0    1
2          2          2          3          1        3  26.0    0
3          3          3          4          1        1  35.0    1
4          4          4          5          0        3  35.0    0

      Parch    Fare  AgeImputed  ...  Embarked_missing  CabinLetter_A  \
0         0   7.2500         False  ...              0              0
1         0  71.2833         False  ...              0              0
2         0   7.9250         False  ...              0              0
3         0  53.1000         False  ...              0              0
4         0   8.0500         False  ...              0              0

      CabinLetter_B  CabinLetter_C  CabinLetter_D  CabinLetter_E  CabinLetter_F  \
0                 0              0              0              0              0
1                 0              1              0              0              0
2                 0              0              0              0              0
3                 0              1              0              0              0
4                 0              0              0              0              0

      CabinLetter_G  CabinLetter_T  CabinLetter_missing
0                 0              0                  1
1                 0              0                  0
2                 0              0                  1
3                 0              0                  0
4                 0              0                  1
```

[5 rows x 28 columns]

3.3 Having a quick look at differences between survived and non-survived passengers

Phew, the data-preprocessing is done! This is often a tedious and time-consuming stage with few ‘endorphin rush’ rewards to be had.

Let’s split our data into survived and non-survived and have a quick look to see anything obvious.

```
[28]: mask = data['Survived'] == 1 # mask for survived passengers
      survived = data[mask]

      # Invert mask (for passengers who died)
      mask = mask == False
      died = data[mask]
```

Now let’s have a quick look at mean values for our two groups. We’ll put them side by side in a new dataframe

```
[29]: summary = pd.DataFrame()
      summary['survived'] = survived.mean()
      summary['died'] = died.mean()
      summary
```

```
[29]:
```

	survived	died
Unnamed: 0	443.368421	446.016393
Unnamed: 0.1	443.368421	446.016393
PassengerId	444.368421	447.016393
Survived	1.000000	0.000000
Pclass	1.950292	2.531876
Age	28.291433	30.028233
SibSp	0.473684	0.553734
Parch	0.464912	0.329690
Fare	48.395408	22.117887
AgeImputed	0.152047	0.227687
EmbarkedImputed	0.005848	0.000000
CabinLetterImputed	0.602339	0.876138
CabinNumberImputed	0.611111	0.885246
male	0.318713	0.852459
Embarked_C	0.271930	0.136612
Embarked_Q	0.087719	0.085610
Embarked_S	0.634503	0.777778
Embarked_missing	0.005848	0.000000
CabinLetter_A	0.020468	0.014572
CabinLetter_B	0.102339	0.021858
CabinLetter_C	0.102339	0.043716
CabinLetter_D	0.073099	0.014572
CabinLetter_E	0.070175	0.014572
CabinLetter_F	0.023392	0.009107

CabinLetter_G	0.005848	0.003643
CabinLetter_T	0.000000	0.001821
CabinLetter_missing	0.602339	0.876138

3.4 Save processed data

```
[30]: data.to_csv('./data/processed_data.csv', index=False)
```